

---

title: "Lab 5"

author: "Feiyi Ma"

output: pdf\_document

date: "11:59PM March 18, 2021"

---

Create a 2x2 matrix with the first column 1's and the next column iid normals. Find the absolute value of the angle (in degrees, not radians) between the two columns.

```
```{r}
```

```
x = matrix(1:1, nrow =2, ncol=2)
```

```
x[,2]=rnorm(2)
```

```
norm_vec = function (v){
```

```
  sqrt(sum(v^2))
```

```
}
```

```
#x[,2]%*%t(x[,1])
```

```
cos_theta = t(x[,1])%*%x[,2]/(norm_vec(x[,1])%*%norm_vec(x[,2]))
```

```
cos_theta
```

```
acos(cos_theta)*180/pi
```

```
```
```

Repeat this exercise `Nsim = 1e5` times and report the average absolute angle.

```
```{r}
```

```
Nsim = 1e5 #100,000
```

```

angles = array(NA,Nsim)
for(i in 1:Nsim) {
  x = matrix(1:1, nrow =2, ncol=2)
  x[,2]=rnorm(2)
  cos_theta = t(x[,1])%*%x[,2]/(norm_vec(x[,1])%*%norm_vec(x[,2]))
  angles[i]=abs(90-acos(cos_theta)*180/pi)
}
mean(angles)
...

```

Create a nx2 matrix with the first column 1's and the next column iid normals. Find the absolute value of the angle (in degrees, not radians) between the two columns. For n = 10, 50, 100, 200, 500, 1000, report the average absolute angle over `Nsim = 1e5` simulations.

```

```{r}
#return 0-360 or -180 to 180
N=c(2,5,10,50,100, 200, 500, 1000)
Nsim = 1e5
angles = matrix(NA,nrow=Nsim,ncol =length(N))
for(j in 1:length(N)){
  for(i in 1:Nsim) {
    x = matrix(1, nrow =N[j], ncol=2)
    x[,2]=rnorm(N[j])
    cos_theta = t(x[,1])%*%x[,2]/(norm_vec(x[,1])%*%norm_vec(x[,2]))
    angles[i,j]=abs(90-acos(cos_theta)*180/pi)
  }
}
colMeans(angles) #1000 damnation is 1.44 degree from orthogonal
...

```

What is this absolute angle converging to? Why does this make sense?

absolute angles form 90 to it converging to 0. This make sense because in high Dim orthogonal.

Create a vector  $y$  by simulating  $n = 100$  standard iid normals. Create a matrix of size  $100 \times 2$  and populate the first column by all ones (for the intercept) and the second column by 100 standard iid normals. Find the  $R^2$  of an OLS regression of  $y \sim X$ . Use matrix algebra.

```
```{r}
n=100
X = cbind(1,rnorm(n))
y = rnorm(n)
H = X%*%solve(t(X)%*%X)%*% t(X)
y_hat = H %*%y
y_bar = mean(y)
SSR =sum((y_hat-y_bar)^2)
SST = sum((y-y_bar)^2)

Rsqr = (SSR/SST)
Rsqr
```
```

Write a for loop to each time bind a new column of 100 standard iid normals to the matrix  $X$  and find the  $R^2$  each time until the number of columns is 100. Create a vector to save all  $R^2$ 's. What happened??

```
```{r}

#need to run the above function in order to run this one

Rsqr_s = array(NA,dim = n-2)
```

```

for (j in 1:(n-2)){
  X=cbind(X,rnorm(n))
  H = X%%solve(t(X)%%X)%% t(X)
  y_hat = H %% y
  y_bar = mean(y)
  SSR =sum((y_hat-y_bar)^2)
  SST = sum((y-y_bar)^2)

  Rsq_s[j] = (SSR/SST)
}
Rsq_s
diff(Rsq_s)
...

```

Test that the projection matrix onto this X is the same as  $I_n$ . You may have to vectorize the matrices in the `expect\_equal` function for the test to work.

```

``{r}
pacman::p_load(testthat)
dim(x) # y_hat = h_hat
H = X%%solve(t(X)%%X)%% t(X) #identity matrix
H[1:10,1:10] # 0 for all of them
I = diag(n)
expect_equal(I,H) # it passed after 10^(-7) or 10^(-8)
...

```

Add one final column to X to bring the number of columns to 101. Then try to compute  $R^2$ . What happens?

```

```{r}
X=cbind(X,rnorm(n))
H = X%%solve(t(X)%%X)%% t(X)
#solve will fail by 1 column,can invert the 101 by 101 in rank 100
y_hat = H %%y
y_bar = mean(y)
SSR =sum((y_hat-y_bar)^2)
SST = sum((y-y_bar)^2)

rsq = (SSR/SST)
rsq
```

```

Why does this make sense?

x<sub>-inverse</sub> & x can not invert due to rank different. solve will fail by 1 column,can invert the 101 by 101 in rank 100

Write a function spec'd as follows:

```

```{r}
#' Orthogonal Projection
#'
#' Projects vector a onto v.
#'
#' @param a the vector to project
#' @param v the vector projected onto
#'
#' @returns a list of two vectors, the orthogonal projection parallel to v named a_parallel,

```

#' and the orthogonal error orthogonal to v called a\_perpendicular

```
orthogonal_projection = function(a, v){  
  H = v%*%t(v) / norm_vec(v)^2 #dimension is  
  a_parallel = H%*%a  
  a_perpendicular = a - a_parallel
```

```
  list(a_parallel = a_parallel, a_perpendicular = a_perpendicular)  
}  
...
```

Provide predictions for each of these computations and then run them to make sure you're correct.

```
```${r}  
orthogonal_projection(c(1,2,3,4), c(1,2,3,4))  
#prediction:  
orthogonal_projection(c(1, 2, 3, 4), c(0, 2, 0, -1))  
#prediction:  
result = orthogonal_projection(c(2, 6, 7, 3), c(1, 3, 5, 7))  
# v = c(1, 3, 5, 7) the direction of projection  
t(result$a_parallel) %*% result$a_perpendicular  
#prediction:  
result$a_parallel + result$a_perpendicular  
#prediction: give the original vector  
result$a_parallel / c(1, 3, 5, 7)  
#prediction: % of the orthogonal line present from the original line  
...
```

Let's use the Boston Housing Data for the following exercises

```
```{r}
```

```
#model to be worked on
```

```
y = MASS::Boston$medv
```

```
X = model.matrix(medv ~ ., MASS::Boston)
```

```
p_plus_one = ncol(X)
```

```
n = nrow(X)
```

```
head(X)
```

```
```
```

Using your function `orthogonal_projection` orthogonally project onto the column space of X by projecting y on each vector of X individually and adding up the projections and call the sum `yhat_naive`.

```
```{r}
```

```
yhat_naive = rep(0,n)
```

```
for(j in 1:p_plus_one){ #col use j
```

```
  yhat_naive = yhat_naive + orthogonal_projection(y,X[,j])$a_parallel
```

```
  # like += itself  second is a matrix need to pull the data out
```

```
}
```

```
```
```

How much double counting occurred? Measure the magnitude relative to the true LS orthogonal projection.

```

```{r}
yhat = X %*% solve(t(X) %*% X) %*% t(X) %*% y
sqrt(sum(yhat_naive^2)) / sqrt(sum(yhat^2))
#yhat_naive not y_hat, many double double counting
```

```

Is this ratio expected? Why or why not?

#TO-DO yhat\_naive not y\_hat, many double double counting. It's expected to be different from 1.

Convert X into V where V has the same column space as X but has orthogonal columns. You can use the function `orthogonal\_projection`. This is the Gram-Schmidt orthogonalization algorithm.

```

```{r}
V = matrix(NA, nrow = n, ncol = p_plus_one)
V[, 1] = X[, 1]
for (j in 2:p_plus_one){
  V[,j] = X[,j]
  for (k in 1:(j-1)){
    V[,j] = V[,j] - orthogonal_projection(X[,j], V[,k])$a_parallel
  }
}
head(X)
V[,7] %*% V[,9]
```

```

Convert V into Q whose columns are the same except normalized

```

```{r}

```



```

Q = matrix(NA, nrow = n, ncol = p_plus_one)
for (j in 1:p_plus_one){
  Q[,j] = V[, j] / norm_vec(V[,j])
}
#TO-DO
'''

```

Verify  $Q^T Q$  is  $I_{p+1}$  i.e.  $Q$  is an orthonormal matrix.

```

'''{r}
expect_equal(t(Q)%*%Q, diag(p_plus_one))
'''

```

Is your  $Q$  the same as what results from R's built-in QR-decomposition function?

```

'''{r}
Q_from_Rs_builtin = qr.Q(qr(x))
#qr.Q is the get function
expect_equal(Q,Q_from_Rs_builtin) # 7084 = 506*18 matrix,
'''

```

Is this expected? Why did this happen?

#To-do Yes, the  $Q$  matrix is 506 by 18 = 7084, the original matrix, the qr of  $Q$  matrix is the orthogonal projection of the matrix. It shows only part of the matrix is orthogonal to the original one.

Project  $y$  onto  $\text{colsp}[Q]$  and verify it is the same as the OLS fit. You may have to use the function ``unnname`` to compare the vectors since they the entries will likely have different names.

```

```{r}
yhat_naive = rep(0,ncol(Q))
for(j in 1:ncol(Q)){ #col use j

  yhat_naive = yhat_naive + orthogonal_projection(y,Q[,j])$a_parallel
  # like += itself second is a matrix need to pull the data out

}
expect_equal(yhat_naive,orthogonal_projection(y,Q)$a_parallel)
```

```

Project  $y$  onto  $\text{colsp}[Q]$  one by one and verify it sums to be the projection onto the whole space.

```

```{r}
yhat_naive = rep(0,ncol(Q))
for(j in 1:ncol(Q)){ #col use j

  yhat_naive = yhat_naive + orthogonal_projection(y,Q[,j])$a_parallel
  # like += itself second is a matrix need to pull the data out

}
expect_equal(sum(yhat_naive),sum(orthogonal_projection(y,Q)$a_parallel))
```

```

Split the Boston Housing Data into a training set and a test set where the training set is 80% of the observations. Do so at random.

```

```{r}
K = 5
n_test= round(n* 1/K)

```

```

n_train = n- n_test
test_indices = sample(1:n, n * 1 / K)
train_indices = setdiff(1:n , test_indices)
y=MASS::Boston$medv
y_bar = mean(y)
SST = sum((y-y_bar)^2)
x = as.matrix((cbind(1,MASS::Boston[,1:13])))
n = nrow(x)

x_train = x[train_indices,]
x_test = x[test_indices,]
y_train = y[train_indices]
y_test = y[test_indices]

dim(x_train)
dim(x_test)
length(y_train)
length(y_test)

...

```

Fit an OLS model. Find the  $s_e$  in sample and out of sample. Which one is greater? Note: we are now using  $s_e$  and not RMSE since RMSE has the  $n-(p + 1)$  in the denominator not  $n-1$  which attempts to de-bias the error estimate by inflating the estimate when overfitting in high  $p$ . Again, we're just using  $\text{sd}(e)$ , the sample standard deviation of the residuals.

```

```{r}
mod = lm(y_train ~ ., data.frame(x_train))
summary(mod)$r.squared
sd(mod$residuals)

```

```
...
```

Do these two exercises `Nsim = 1000` times and find the average difference between  $s_e$  and  $ooss_e$ .

```
``{r}
```

```
K= 5
```

```
n_test= round(n* 1/K)
```

```
n_train = n- n_test
```

```
Nsim = 1000
```

```
ooss_e = array(NA, dim = Nsim)
```

```
s_e = array(NA, dim = Nsim)
```

```
for (i in 1:Nsim){
```

```
  test_indices = sample(1:n, 1 / K * n)
```

```
  train_indices = setdiff(1:n , test_indices)
```

```
  #y=MASS::Boston$medv
```

```
  #y_bar = mean(y)
```

```
  #SST = sum((y-y_bar)^2)
```

```
  #x = as.matrix((cbind(1,MASS::Boston[,1:13])))
```

```
  #n = nrow(x)
```

```
  x_train = x[train_indices,]
```

```
  x_test = x[test_indices,]
```

```
  y_train = y[train_indices]
```

```
  y_test = y[test_indices]
```

```
  N_matrix = lm(y_train~.+0, data.frame(x_train))
```

```
  y_hat_oos = predict(N_matrix, data.frame(x_test))
```

```

ooss_e[i] = sd(y_test - y_hat_oos)
s_e[i] = sd(N_matrix$residuals)
}
abs(mean(s_e - ooss_e))

'''

```

We'll now add random junk to the data so that `p_plus_one = n_train` and create a new data matrix `X_with_junk`.

```

'''{r}
#p_plus_one = n_train
x_with_junk = cbind(x, matrix(rnorm(n * (n_train - p_plus_one)), nrow = n))
dim(x)
dim(x_with_junk)
'''

```

Repeat the exercise above measuring the average `s_e` and `ooss_e` but this time record these metrics by number of features used. That is, do it for the first column of `X_with_junk` (the intercept column), then do it for the first and second columns, then the first three columns, etc until you do it for all columns of `X_with_junk`. Save these in `s_e_by_p` and `ooss_e_by_p`.

```

'''{r}
K= 5
n_test= round(n* 1/K)
n_train = n- n_test
ooss_e_by_p = array(NA, dim = ncol(x_with_junk))
s_e_by_p = array(NA, dim = ncol(x_with_junk))
Nsim = 10

```

```

for (i in 1:ncol(x_with_junk)){
  ooss_e = array(NA, dim = Nsim)
  s_e = array(NA, dim = Nsim)
  for(j in 1:Nsim){

test_indices = sample(1:n, 1 / K * n)
train_indices = setdiff(1:n , test_indices)

x_train = x_with_junk[train_indices, 1:j, drop= FALSE]
x_test = x_with_junk[test_indices, 1:j, drop= FALSE]
y_train = y[train_indices]
y_test = y[test_indices]

N_matrix = lm(y_train~.+0, data.frame(x_train))
y_hat_oos = predict(N_matrix, data.frame(x_test))
ooss_e[i] = sd(y_test - y_hat_oos)
s_e[i] = sd(N_matrix$residuals)
  }
  ooss_e_by_p[j] = mean(ooss_e)
  s_e_by_p[j] = mean(s_e)
}
ooss_e_by_p
s_e_by_p
...

```

You can graph them here:

```

``{r}
pacman::p_load(ggplot2)
ggplot(
  rbind(
    data.frame(s_e = s_e_by_p, p = 1 : n_train, series = "in-sample"),
    data.frame(s_e = ooss_e_by_p, p = 1 : n_train, series = "out-of-sample")
  )) +
  geom_line(aes(x = p, y = s_e, col = series))
``

```

Is this shape expected? Explain.

#TO-DO Should be, as we add more junk, there are small part orthogonal to the regression as overfitting all the regression.